

SUMMARY STATISTICS

Applied math for data science (DSC 482/MTH 445) Fall 2022

Table of Contents

- [1. Measures of centrality](#)
- [2. Preparations to code along](#)
- [3. Mean or arithmetic average](#)
- [4. Median or middle magnitude](#)
- [5. Mode or most common observation](#)
- [6. Mean and median with built-in functions](#)
- [7. Mode with contingency table](#)
- [8. Practice: measures of centrality](#)
 - [8.1. SOLUTION](#)
- [9. Missing or undefined values](#)
- [10. Practice: missing values](#)
 - [10.1. Solution](#)
- [11. excuRsion: category subsets with tapply \(Matloff\)](#)
- [12. Practice: mean weight of chicks by feed type](#)
- [13. Counting chicks](#)
- [14. Visualizing contingency tables](#)
 - [14.1. Solution](#)
- [15. Proportions with table](#)
- [16. Proportions with logical flag vectors](#)
- [17. Proportions with tapply](#)
- [18. Rounding numeric data with round](#)
- [19. Percentages vs. proportion](#)
- [20. Exercises \(for home\)](#)
- [21. TODO Glossary: concepts](#)
- [22. TODO Glossary: code](#)
- [23. References](#)



Figure 1: poster for "The Terminal List" (Amazon, 2022-).

- Centrality: Mean, Median, Mode
- Counts, Percentages, and Proportions
- Quantiles, Percentiles, and 5-number-summary
- Spread: Variance, Standard Deviation and Interquartile Range
- 2 Practice exercises and 1 exercise session

1 Measures of centrality



- Explain large collections of data
- Describe where numeric observations are centered
- Centrality measures suggest symmetries where none exist

2 Preparations to code along

```

#+PROPERTY: header-args:R :export both :results output :session *R*

#+begin_src R
  head(mtcars)
#+end_src

#+RESULTS:
:
:      mpg  cyl  disp  hp drat   wt  qsec vs am gear carb
: Mazda RX4      21.0   6  160 110 3.90 2.620 16.46 0  1    4    4
: Mazda RX4 Wag  21.0   6  160 110 3.90 2.875 17.02 0  1    4    4
: Datsun 710     22.8   4  108  93 3.85 2.320 18.61 1  1    4    1
: Hornet 4 Drive  21.4   6  258 110 3.08 3.215 19.44 1  0    3    1
: Hornet Sportabout 18.7   8  360 175 3.15 3.440 17.02 0  0    3    2
: Valiant        18.1   6  225 105 2.76 3.460 20.22 1  0    3    1

-\\***- plot.org All L1 (Org)

```

- **Open** a new Org-mode file stats.org in Emacs
- **Add** this line at the top of the file plot.org:


```
#+PROPERTY: header-args:R :results output :session *R*
```
- **Activate** the code by putting your cursor on the line and entering C-c C-c. You should see the message Local setup has been refreshed in the minibuffer at the bottom of the editor.
- When you execute your first R code block, you'll be asked where you want the session named *R* to run: **enter** the path to plot.org
- For plots, **use** the header :results graphics file :file plot.png
- When you leave Emacs, you'll be warned that the session *R* is active: you can **ignore** this warning

3 Mean or arithmetic average

- For a set of n labeled numeric measurements, the sample mean is the arithmetic average over all

$$\bar{x} = \frac{(x_1 + x_2 + \dots + x_n)}{n} = \frac{1}{n} \sum_{i=1}^n x_i$$

measurements:

- If you observe 8 points 2, 4.4, 3, 3, 2, 2.2, 2, 4, the mean is:

```

x <- c(2,4.4,3,3,2,2.2,2,4) # store observations in vector
x_mean <- sum(x)/length(x) # arithmetic mean
x_mean # print mean

```

4 Median or middle magnitude

- Sort your observations by magnitude
- For an odd number of observations: take the middle value
- For an even number of observations: average two middle values

$$\bar{m}_x = \begin{cases} x_i^{(\frac{n+1}{2})}, & \text{if } n \text{ is odd} \\ \left(x_i^{(\frac{n}{2})} + x_j^{(\frac{n}{2}+1)} \right) / 2, & \text{if } n \text{ is even} \end{cases}$$

Where the upper index denotes the order statistics: $x_i^{(t)}$ is the t -th smallest observation regardless of the observation index i .

- If you observe 8 points 2, 4.4, 3, 3, 2, 2.2, 2, 4, you have $n/2=4$.

```
x <- c(2,4.4,3,3,2,2.2,2,4) # store observations in vector
sorted_x <- sort(x) # sort observations
sorted_x
# median by hand:
(sorted_x[length(x)/2] + sorted_x[length(x)/2+1])/2
```

5 Mode or most common observation

- Used with numeric-discrete data than numeric-continuous
- Used when discussing probability *density* functions
- Collection of numeric measurements may have no or > 1 mode
- If you observe 8 points 2, 4.4, 3, 3, 2, 2.2, 2, 4, you can tabulate the frequency of each measurement:

Observation	2	2.2	3	4	4.4
Frequency	3	1	2	1	1

6 Mean and median with built-in functions

- Create a new Org-mode file

- Create a R code block with header: R :results output :session
- Store the eight observations as a numeric vector xdata:

{2, 4.4, 3, 3, 2, 2.2, 2, 4}

```
xdata <- c(2,4.4,3,3,2,2.2,2,4) # store observations in vector  
xdata
```

- Compute the mean and the median and store them in variables.

```
x.bar <- mean(xdata)  
x.bar
```

```
m.bar <- median(xdata)  
m.bar
```

7 Mode with contingency table

- To find a mode, compute the contingency table for xdata.

```
xtab <- table(xdata)  
xtab
```

- To identify the most frequent values automatically, use range which reports min and max of xdata.

```
sort(xdata) # sort vector values  
min(xdata)  
max(xdata)  
range(xdata) # return min and max value
```


- Applying these functions to a table operates on the frequencies:

```
xtab # object that stores the contingency table  
max(xtab) # max frequency in the table
```

- Finally, use a logical index vector to get the mode:

```
d.bar <- xtab[xtab == max(xtab)]  
d.bar # prints the value and the frequency
```

8 Practice: measures of centrality



1. Calculate the mean and median weights of the chicks in chickwts
2. Identify the value and frequency of the most common earthquake magnitude in quakes

8.1 SOLUTION

1. Calculate the mean and median weights of the chicks in chickwts

```
str(chickwts)
mean(chickwts$weight)
median(chickwts$weight)
```

```
'data.frame': 71 obs. of 2 variables:
 $ weight: num 179 160 136 227 217 168 108 124 143 140 ...
 $ feed : Factor w/ 6 levels "casein","horsebean",...: 2 2 2 2 2 2 2 2 2 2 ...
[1] 261.3099
[1] 258
```

Answer: the average chick weight 261 grams, and the chicken in the middle weighed 258 grams.

2. Identify the value and frequency of the most common earthquake magnitude in quakes

```
str(quakes)
Qtab <- table(quakes$mag)
Qtab[Qtab == max(Qtab)]
```

```
'data.frame': 1000 obs. of 5 variables:
 $ lat : num -20.4 -20.6 -26 -18 -20.4 ...
 $ long : num 182 181 184 182 182 ...
 $ depth : int 562 650 42 626 649 195 82 194 211 622 ...
 $ mag : num 4.8 4.2 5.4 4.1 4 4 4.8 4.4 4.7 4.3 ...
 $ stations: int 41 15 43 19 11 12 43 15 35 19 ...
4.5
107
```


Answer: the most common earthquakes were 107 occurrences of magnitude 4.5.

9 Missing or undefined values

- Many of the standard stats functions in R will not run if the data set contains missing (NA) or undefined (NaN) values.

```
mean(c(1,4,NA))
mean(c(1,4,NaN))
```

- To prevent inclusion of these special values, switch on the `na.rm` attribute.

```
mean( c(1,4,NA), na.rm = TRUE)
mean( c(1,4,NaN), na.rm = TRUE)
```

- Anything that calculates a numeric statistic based on a numeric vector carries this attribute: `sum`, `mean`, `median`, `max`, `min`, `range`.

10 Practice: missing values

- Look at the `Pima.tr` dataset in the `MASS` package.
- Use `summary` to find out how many values are missing in the measurements of the body mass index.
- What is the sample mean of the body mass index values?

10.1 Solution

```
library(MASS) # load MASS package
str(Pima.tr2) # look at structure of Pima.tr2
summary(Pima.tr2) # summarize the stats of the data set
mean(Pima.tr2$bmi) # this mean cannot be computed: NA
mean(Pima.tr2$bmi, na.rm=TRUE) # the 3 NAs have been removed
```

11 excuRsion: category subsets with `tapply` ([Matloff](#))

- The built-in `ToothGrowth` data set contains the numeric variable `len` (length of a tooth), and the categorical variable `supp` with two levels, `OJ` and `VC` for "Orange juice" and "Vitamin C".

```
str(ToothGrowth)
```

- Let's say we want to know the mean length `ToothGrowth$len` for each of the two levels.
- The `tapply` function allows us to split the vector `X = ToothGrowth$len` in two groups according to the values of `INDEX = ToothGrowth$supp`, and then apply the function `FUN = mean`.

```
tapply(X=ToothGrowth$len, INDEX=ToothGrowth$supp, FUN=mean)
```


- [X] Check out `help(tapply)` - remember to enter system commands in the R console (**why is help a system command?**)

12 Practice: mean weight of chicks by feed type

- []

If you want to find the mean weight of chicks grouped by feed type, you could use `mean` on each specific subset - how would this look like?

```
mean(chickwts$weight[chickwts$feed == "casein"])
mean(chickwts$weight[chickwts$feed == "horsebean"])
mean(chickwts$weight[chickwts$feed == "linseed"])
mean(chickwts$weight[chickwts$feed == "meatmeal"])
mean(chickwts$weight[chickwts$feed == "soybean"])
mean(chickwts$weight[chickwts$feed == "sunflower"])
```

- []

Instead, use `tapply(X, INDEX, FUN)` to calculate these values using just one line of code! Remember: `X` is the vector, `INDEX` is the splitting category (factor level), and `FUN` is the function.

```
tapply(
  X = chickwts$weight,
  INDEX = chickwts$feed,
  FUN = mean)
```

13 Counting chicks

- Sometimes it's useful to summarize non-numerical data, e.g. the number of observations that fall in a particular category
- *Counts* or *frequencies* are summary statistics of categorical data
- Again you can use the contingency table command for frequencies - e.g. for the feed types in the `chickwts` data set.

```
table(chickwts$feed)
```

14 Visualizing contingency tables

- [] **How would you visualize this table?** Think about the data and about generic plotting in R.

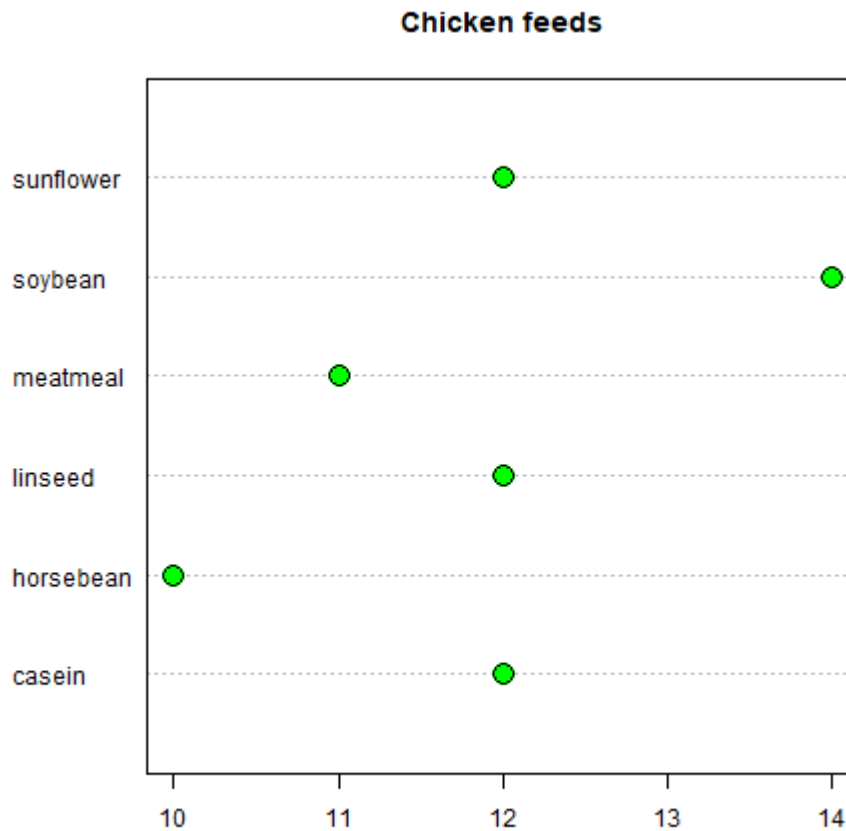
14.1 Solution

- Plug the table into `plot` - it's a generic graphic function and it does have a `plot.table` method.

```
plot(x=table(chickwts$feed),
     main="Chicken feeds",
     ylab="Frequencies")
```

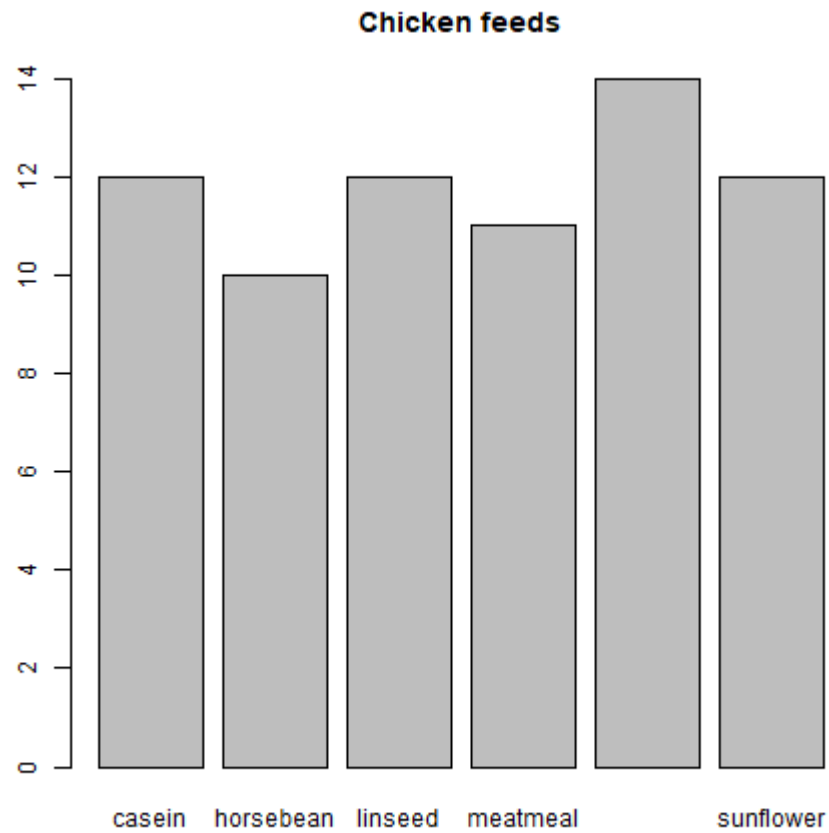
- "Thin frequencies" like this are better represented as dots in a so-called dot plot, with the `dotchart` function (is it generic?):

```
dotchart(x=table(chickwts$feed),  
         main="Chicken feeds",  
         pch=21, bg="green", pt.cex=2)
```

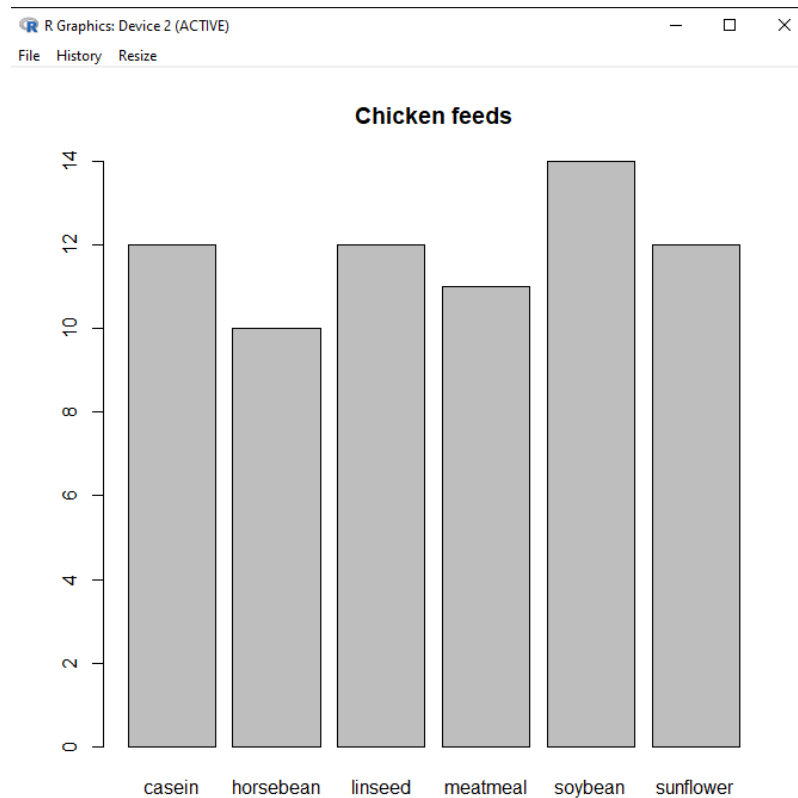


- Since the table contains just categories, you could also try and plug them directly into `plot`, without going through `table`:

```
plot(x=chickwts$feed,  
     main="Chicken feeds")
```



- **Bonus (10 pts):** find out why the last graph does not show the level soybean? When executed in the R console, the command works:



15 Proportions with table

- More information from counts includes the proportion of observations that fall into each category.
- Proportions represent the fraction of observations in each category, usually as a decimal number between 0 and 1, and they add up to 1.
- For numeric vectors like our sample vector, there is a function, `prop.table`.

```
x <- c(2,4.4,3,3,2,2.2,2,4) # store observations in vector
prop.table(x) # prop.table works for numeric vectors
sum(prop.table(x)) # proportions add up to 1
```

- For category vectors, or factors with levels, you need to divide the count by the overall sample size, which is `nrow(chickwts)`.

```
table(chickwts$feed) / nrow(chickwts)
table(chickwts$feed)
sum(table(chickwts$feed)) # counts add up to total no. of records
nrow(chickwts) # number of rows in the data set
```

16 Proportions with logical flag vectors

- You do not always need `table` - the sum over a logical flag vector is just as good because such a vector of TRUE and FALSE is coded as a vector of 0 and 1.
- Example: `chickwts$feed == "soybean"` lists all chicks fed with soybean as TRUE (or 1):

```
chickwts$feed == "soybean"
```

- For example, to find the proportion of chicks fed soybean:

```
sum(chickwts$feed == "soybean") / nrow(chickwts)
```

- This is equivalent to averaging over the logical flag vector:

```
mean(chickwts$feed == "soybean")
```

- You can use this approach to calculate the proportion of entities in groups. E.g. the proportion of chicks fed soybean or horsebean:

```
mean(chickwts$feed == "soybean" | chickwts$feed == "horsebean")
```

- This computation uses the following logical argument vector:

```
chickwts$feed == "soybean" | chickwts$feed == "horsebean"
```

17 Proportions with `tapply`

- You can also use `tapply` with the `FUN` argument to be an anonymous (non-named) function that computes the mean for each feed level

```
prop <- tapply(
  X = chickwts$weight, # object that can be split by factor levels
  INDEX = chickwts$feed, # list of factors
  FUN = function(x) length(x)/nrow(chickwts) # function to be applied
                                           # to factors
  prop
  sum(prop)
```

- Here, the anonymous function is defined with a dummy argument `x`.

18 Rounding numeric data with `round`

- The `round` function rounds numeric data output to a certain number of decimal places. It has only two arguments, input data and digits.

```
round(
  table(chickwts$feed) / nrow(chickwts), # input data
```

```
digits = 3) # output digits
```

19 Percentages vs. proportion

- Percentage and proportion represent the same thing.
- They differ in scale - percentage is proportion multiplied by 100.
- The percentage of chicks on a soybean diet is approximately 19.7%

```
round(  
  x = mean(chickwts$feed == "soybean") * 100,  
  digits = 1)
```

- Proportions always lie in $[0,1]$ while percentages lie in $[0,100]$.
- Statisticians prefer percentages when discussing percentiles, and proportions when discussing probabilities.

20 Exercises (for home)



- Download practice file [from GitHub](#) and save as Org-mode file
- Submit completed Org-mode file [to Canvas](#) by Thursday, 22-Sept, 8 am

21 **TODO** Glossary: concepts

TERM	MEANING
------	---------

22 **TODO** Glossary: code

CODE	MEANING
------	---------

23 References

- DataCamp (n.d.). Introduction to Statistics. URL: datacamp.com.
- Davies TD (2016). Book of R. NoStarch Press. URL: nostarch.com
- Matloff N (2022). fasteR. URL: github.com/matloff/fasteR

Author: MARCUS BIRKENKRAHE

Created: 2022-09-22 Thu 17:59